

MALVIKA SANGHVI

Seattle, WA | 206-877-2948 | malvikasanghvi@gmail.com | linkedin.com/malvikasanghvi

SUMMARY

AI Product Manager with 7 years across logistics and real estate who builds AI products end to end and grounds them in evaluation so teams can trust the output. Recently built a multi-LLM assistant at Radiant Global Logistics, and independently shipped SkinScore AI (verifies product claims against ingredient evidence) and Isaiah, a voice agent that helps older adults troubleshoot technology by phone.

SKILLS

- **AI & ML:** LangChain, Google ADK, Multi-Agent Orchestration, Multi-MCP Architecture, RAG, Function Calling, Prompt Engineering, LangSmith, Anthropic API (Claude), OpenAI API, Azure OpenAI, Claude Code, Cursor, TensorFlow, Scikit-learn
- **Analytics & Experimentation:** SQL, Python, Google Analytics, Snowflake, Tableau, Power BI, Amplitude, Mixpanel, Heap, Segment, Optimizely, LaunchDarkly
- **Cloud & Infrastructure:** GCP (Cloud Run), AWS, Azure, Azure AI Foundry, SAP S/4HANA, Postgres, PGVector, Zapier, n8n
- **CRM & Marketing Automation:** Salesforce, HubSpot, Marketo, Apollo, Clay
- **Product & Collaboration:** Jira, Confluence, Figma, Miro, Notion, Slack, Asana

EDUCATION

University of Washington – Foster School of Business | Seattle, WA

Mar 2025 – Jun 2026

M.S. Information Systems · GPA: 3.76

Georgia Institute of Technology | Atlanta, GA

Aug 2012 – Dec 2016

B.S. Biomedical Engineering

WORK EXPERIENCE

Radiant Global Logistics

Seattle, WA

AI Product Manager (Intern)

Mar 2026 – Jun 2026

Publicly-traded North American 3PL (NYSE: RLGT, \$900M+ revenue) advancing AI-driven supply chain transformation.

- Built a multi-LLM AI agent platform on GCP with Google ADK and 5 MCP servers, routing across Claude, GPT, and Gemini by task type, grounded in operational data via RAG (PGVector), then validated tool accuracy and workflow fit with 200+ cross-functional teams across operations, sales, customer service, and finance, using their feedback and eval-pipeline reliability data to close the gaps that blocked real workflows.
- Built online and offline evaluation, an LLM-as-judge scoring responses across 7 quality dimensions plus deterministic runtime assertions, then extended it to multi-agent trace analysis to catch sequencing failures invisible at the single-call level, reaching 0.96 faithfulness, 0.94 groundedness, and a 94% assertion pass rate.
- Built responsible AI in from the start: grounded responses in governed data to prevent fabrication, and kept sensitive operational data inside governed infrastructure.
- Built query-based routing that sent the highest-volume shipper and station-manager questions to a lightweight fast path and reserved the full multi-agent pipeline for the long tail, cutting cost and latency while holding quality.
- Made the call to route tools through MCP rather than direct native tool calls when engineering pushed back on the added latency and the load of maintaining separate servers, arguing it on reuse across systems instead of per-system wiring, and carried the decision as the platform grew to 5 MCP servers.

Pratham Realty

Mumbai, India

Lead Product Manager

Aug 2021 – Jun 2024

Leading Indian real estate developer with 10M+ sq ft of completed and sold projects.

- Owned a company-wide SAP S/4HANA implementation across 10 teams whose metrics fed into each other's formulas under different definitions, mapping those dependencies into one shared data model and governance standard, and making the custom-versus-standard call on each gap on ROI rather than preference, cutting compliance flags per audit cycle from 48% to 21%.
- Led product roadmap for automated workflows and executive dashboards across 14 teams, partnering with engineering on requirements that cut cycle time 80% and increased retention 65%.

Product Manager

Jan 2017 – Jul 2021

- Drove \$2.5M revenue and a 27% qualified-lead increase through data-driven repositioning and analytics instrumentation across digital, broker, and event channels.

INDEPENDENT AI PROJECTS

SkinScoreAI

Feb 2026 – present

- Designed and built SkinScore AI to verify product claims against real ingredient evidence and flag unsupported ones, keeping the scoring engine fully deterministic with LLMs only at the edges so identical inputs return identical, regression-testable verdicts and safety blocks fire every time.
- Built the supporting stack, hybrid RAG retrieval over pgvector, an ingredient knowledge base, and an eval harness with LLM-as-judge grading and multi-agent adversarial review.

Isaiah – Voice AI Agent

May 2026

- Built an inbound voice AI agent on ElevenLabs and Twilio with Azure AI Foundry, with RAG-grounded guardrails and Cal.com tool calling for autonomous human escalation.